





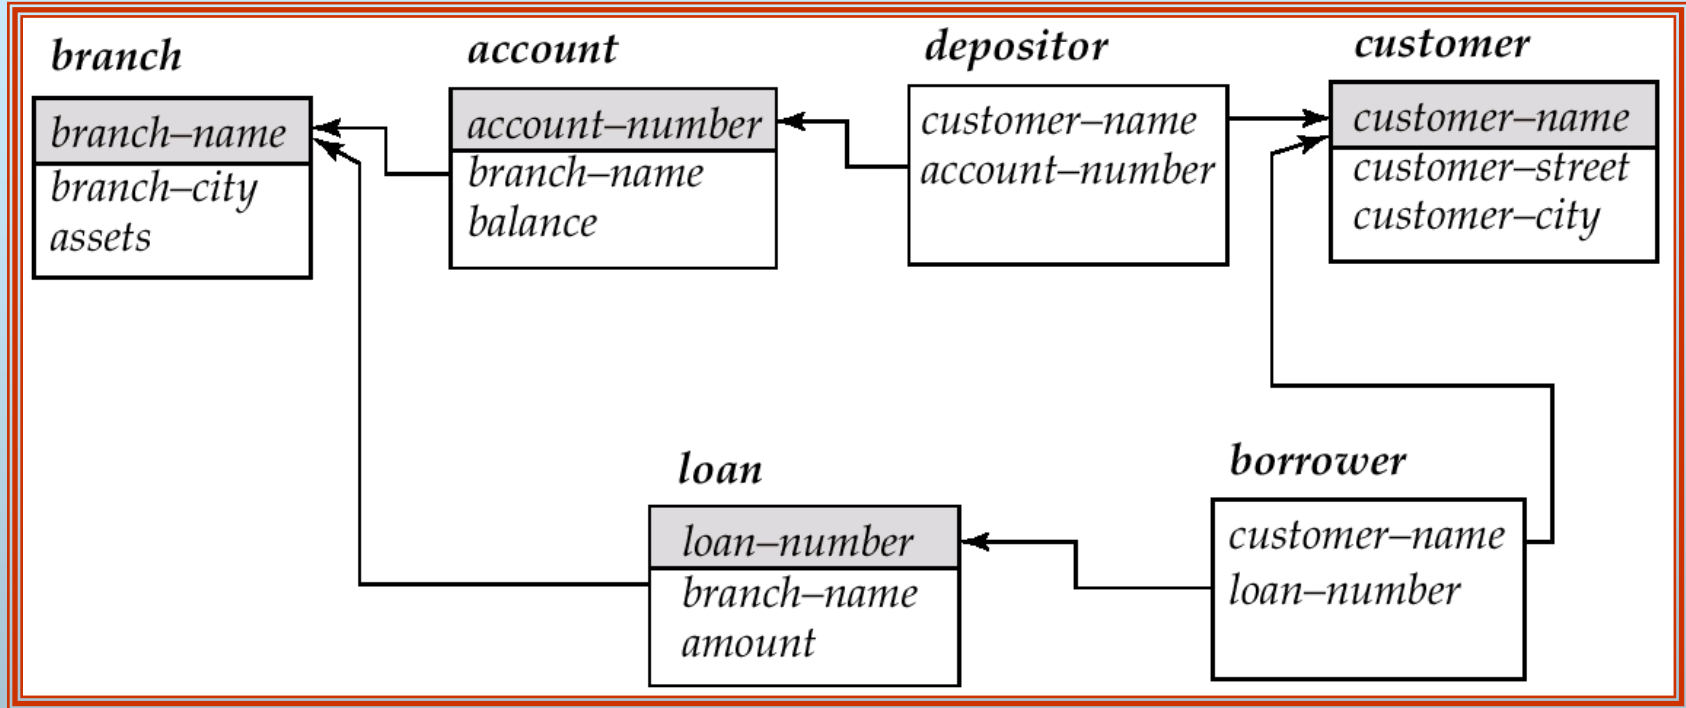
# Chapter 4: SQL

- Basic Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Derived Relations
- Views
- Modification of the Database
- Joined Relations
- Data Definition Languages





# Schema Used in Examples





# Basic Structure

- SQL is based on set and relational operations with certain modifications and enhancements

- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

☞  $A_i$ s represent attributes


☞  $r_i$ s represent relations

☞  $P$  is a predicate.

- This query is equivalent to the relational algebra expression.

$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

- The result of an SQL query is a relation.





# The select Clause

- The **select** clause list the attributes desired in the result of a query

☞ corresponds to the projection operation of the relational algebra

- E.g. find the names of all branches in the *loan* relation

**select** *branch-name*  
**from** *loan*

- In the “pure” relational algebra syntax, the query would be:

$\Pi_{\text{branch-name}}(\text{loan})$


- **NOTE:** SQL does not permit the ‘-’ character in names,

☞ Use, e.g., *branch\_name* instead of *branch-name* in a real implementation.

☞ We use ‘-’ since it looks nicer!

- **NOTE:** SQL names are case insensitive, i.e. you can use capital or small letters.

☞ You may wish to use upper case where-ever we use bold font.






# The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after **select**.
- Find the names of all branches in the *loan* relations, and remove duplicates

```
select distinct branch-name  
from loan
```

- The keyword **all** specifies that duplicates not be removed.

```
select all branch-name  
from loan
```





# The select Clause (Cont.)


- An asterisk in the select clause denotes “all attributes”

```
select *  
from loan
```

- The **select** clause can contain arithmetic expressions involving the operation, +, −, \*, and /, and operating on constants or attributes of tuples.
- The query:

```
select loan-number, branch-name, amount * 100  
from loan
```


would return a relation which is the same as the *loan* relations, except that the attribute *amount* is multiplied by 100.





# The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - ☞ corresponds to the selection predicate of the relational algebra.
- To find all loan number for loans made at the Perryridge branch with loan amounts greater than \$1200.  
**select** *loan-number*  
**from** *loan*  
**where** *branch-name* = 'Perryridge' **and** *amount* > 1200
- Comparison results can be combined using the logical connectives **and**, **or**, and **not**.
- Comparisons can be applied to results of arithmetic expressions.






# The where Clause (Cont.)

- SQL includes a **between** comparison operator
- E.g. Find the loan number of those loans with loan amounts between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )

```
select loan-number
      from loan
      where amount between 90000 and 100000
```








# The from Clause

- The **from** clause lists the relations involved in the query
  - ☞ corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *borrower x loan*  
**select** \*  
**from** *borrower, loan*
- Find the name, loan number and loan amount of all customers having a loan at the Perryridge branch.

```
select customer-name, borrower.loan-number, amount  
from borrower, loan  
where borrower.loan-number = loan.loan-number and  
branch-name = 'Perryridge'
```






# The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Find the name, loan number and loan amount of all customers; rename the column name *loan-number* as *loan-id*.

```
select customer-name, borrower.loan-number as loan-id, amount  
from borrower, loan  
where borrower.loan-number = loan.loan-number
```





# Tuple Variables

- Tuple variables are defined in the **from** clause via the use of the **as** clause.
- Find the customer names and their loan numbers for all customers having a loan at some branch.

```
select customer-name, T.loan-number, S.amount  
from borrower as T, loan as S  
where T.loan-number = S.loan-number
```

- Find the names of all branches that have greater assets than some branch located in Brooklyn.

```
select distinct T.branch-name  
from branch as T, branch as S  
where T.assets > S.assets and S.branch-city = 'Brooklyn'
```





# String Operations


- SQL includes a string-matching operator for comparisons on character strings. Patterns are described using two special characters:
  - ☞ percent (%). The % character matches any substring.
  - ☞ underscore (\_). The \_ character matches any character.
- Find the names of all customers whose street includes the substring “Main”.

```
select customer-name  
from customer  
where customer-street like '%Main%'
```

- Match the name “Main%”

```
like 'Main\%' escape '\'
```

- SQL supports a variety of string operations such as
  - ☞ concatenation (using “||”)
  - ☞ converting from upper to lower case (and vice versa)
  - ☞ finding string length, extracting substrings, etc.






# Ordering the Display of Tuples

- List in alphabetic order the names of all customers having a loan in Perryridge branch

```
select distinct customer-name  
from    borrower, loan  
where borrower loan-number = loan.loan-number and  
        branch-name = 'Perryridge'  
order by customer-name
```


- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - 👉 E.g. **order by** *customer-name desc*





# Duplicates

- In relations with duplicates, SQL can define how many copies of tuples appear in the result.
- *Multiset* versions of some of the relational algebra operators – given multiset relations  $r_1$  and  $r_2$ :
  1.  $\sigma_{\theta}(r_1)$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$ , and  $t_1$  satisfies selections  $\sigma_{\theta}$ , then there are  $c_1$  copies of  $t_1$  in  $\sigma_{\theta}(r_1)$ .
  2.  $\Pi_A(r_1)$ : For each copy of tuple  $t_1$  in  $r_1$ , there is a copy of tuple  $\Pi_A(t_1)$  in  $\Pi_A(r_1)$  where  $\Pi_A(t_1)$  denotes the projection of the single tuple  $t_1$ .
  3.  $r_1 \times r_2$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$  and  $c_2$  copies of tuple  $t_2$  in  $r_2$ , there are  $c_1 \times c_2$  copies of the tuple  $t_1 \cdot t_2$  in  $r_1 \times r_2$





# Duplicates (Cont.)

- Example: Suppose multiset relations  $r_1 (A, B)$  and  $r_2 (C)$  are as follows:

$$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$


- Then  $\Pi_B(r_1)$  would be  $\{(a), (a)\}$ , while  $\Pi_B(r_1) \times r_2$  would be  $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$

- SQL duplicate semantics:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

is equivalent to the *multiset* version of the expression:

$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$






# Set Operations

- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations  $\cup$ ,  $\cap$ ,  $-$ .
- Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:

- 👉  $m + n$  times in  $r$  **union all**  $s$
- 👉  $\min(m, n)$  times in  $r$  **intersect all**  $s$
- 👉  $\max(0, m - n)$  times in  $r$  **except all**  $s$







# Set Operations

- Find all customers who have a loan, an account, or both:


```
(select customer-name from depositor)  
union  
(select customer-name from borrower)
```

- Find all customers who have both a loan and an account.

```
(select customer-name from depositor)  
intersect  
(select customer-name from borrower)
```

- Find all customers who have an account but no loan.

```
(select customer-name from depositor)  
except  
(select customer-name from borrower)
```





# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value


**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values





# Aggregate Functions (Cont.)

- Find the average account balance at the Perryridge branch.


```
select avg (balance)  
  from account  
  where branch-name = 'Perryridge'
```

- Find the number of tuples in the *customer* relation.

```
select count (*)  
  from customer
```

- Find the number of depositors in the bank.

```
select count (distinct customer-name)  
  from depositor
```






# Aggregate Functions – Group By

- Find the number of depositors for each branch.

```
select branch-name, count (distinct customer-name)  
  from depositor, account  
  where depositor.account-number = account.account-number  
 group by branch-name
```

**Note:** Attributes in **select** clause outside of aggregate functions must appear in **group by** list






# Aggregate Functions – Having Clause

- Find the names of all branches where the average account balance is more than \$1,200.

```
select branch-name, avg (balance)  
      from account  
      group by branch-name  
      having avg (balance) > 1200
```

**Note:** predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups






# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The predicate **is null** can be used to check for null values.
  - 👉 E.g. Find all loan number which appear in the *loan* relation with null values for *amount*.

```
select loan-number
from loan
where amount is null
```


- The result of any arithmetic expression involving *null* is *null*
  - 👉 E.g.  $5 + \text{null}$  returns null
- However, aggregate functions simply ignore nulls
  - 👉 more on this shortly





# Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*
  - 👉 E.g.  $5 < \text{null}$  or  $\text{null} <> \text{null}$  or  $\text{null} = \text{null}$
- Three-valued logic using the truth value *unknown*:
  - 👉 OR:  $(\text{unknown} \text{ or } \text{true}) = \text{true}$ ,  $(\text{unknown} \text{ or } \text{false}) = \text{unknown}$   
 $(\text{unknown} \text{ or } \text{unknown}) = \text{unknown}$
  - 👉 AND:  $(\text{true} \text{ and } \text{unknown}) = \text{unknown}$ ,  $(\text{false} \text{ and } \text{unknown}) = \text{false}$ ,  
 $(\text{unknown} \text{ and } \text{unknown}) = \text{unknown}$
  - 👉 NOT:  $(\text{not unknown}) = \text{unknown}$
  - 👉 “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*






# Null Values and Aggregates

- Total all loan amounts

```
select sum (amount)  
from loan
```

- ☞ Above statement ignores null amounts
- ☞ result is null if there is no non-null amount
- ☞ All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes.








# Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries.
- A subquery is a **select-from-where** expression that is nested within another query.
- A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality.






# Example Query

- Find all customers who have both an account and a loan at the bank.

```
select distinct customer-name  
from borrower  
where customer-name in (select customer-name  
                        from depositor)
```

- Find all customers who have a loan at the bank but do not have an account at the bank

```
select distinct customer-name  
from borrower  
where customer-name not in (select customer-name  
                        from depositor)
```





# Example Query

- Find all customers who have both an account and a loan at the Perryridge branch

```
select distinct customer-name
from borrower, loan
where borrower.loan-number = loan.loan-number and
       branch-name = "Perryridge" and
       (branch-name, customer-name) in
           (select branch-name, customer-name
            from depositor, account
            where depositor.account-number =
                account.account-number)
```

- **Note:** Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features.

(Schema used in this example)





# Set Comparison

- Find all branches that have greater assets than some branch located in Brooklyn.

```
select distinct T.branch-name  
from branch as T, branch as S  
where T.assets > S.assets and  
S.branch-city = 'Brooklyn'
```

- Same query using > **some** clause

```
select branch-name  
from branch  
where assets > some  
(select assets  
from branch  
where branch-city = 'Brooklyn')
```





# Definition of Some Clause

- $F \text{ <comp> some } r \Leftrightarrow \exists t \in r \text{ s.t. } (F \text{ <comp> } t)$

Where <comp> can be: <, ≤, >, =, ≠

$$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$$

(read: 5 < some tuple in the relation)


$$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$$

$$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \not\equiv \text{not in}$





# Definition of all Clause

■  $F <\text{comp}> \mathbf{all} \ r \Leftrightarrow \forall t \in r (F <\text{comp}> t)$

$(5 < \mathbf{all} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$


$(5 < \mathbf{all} \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \mathbf{all} \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \mathbf{all} \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \mathbf{all}) \equiv \mathbf{not\ in}$

However,  $(= \mathbf{all}) \neq \mathbf{in}$






# Example Query

- Find the names of all branches that have greater assets than all branches located in Brooklyn.


```
select branch-name
from branch
where assets > all
      (select assets
from branch
where branch-city = 'Brooklyn')
```





# Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$








# Example Query

- Find all customers who have an account at all branches located in Brooklyn.

```
select distinct S.customer-name
from depositor as S
where not exists (
    (select branch-name
from branch
where branch-city = 'Brooklyn')
except
    (select R.branch-name
from depositor as T, account as R
where T.account-number = R.account-number and
        S.customer-name = T.customer-name))
```

- (Schema used in this example)
- Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note: Cannot write this query using = **all** and its variants






# Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- Find all customers who have at most one account at the Perryridge branch.

```
select T.customer-name  
from depositor as T  
where unique (
```

```
    select R.customer-name  
    from account, depositor as R  
    where T.customer-name = R.customer-name and  
          R.account-number = account.account-number and  
          account.branch-name = 'Perryridge')
```

- (Schema used in this example)






# Example Query

- Find all customers who have at least two accounts at the Perryridge branch.

```
select distinct T.customer-name
from depositor T
where not unique (
    select R.customer-name
    from account, depositor as R
    where T.customer-name = R.customer-name
and
    R.account-number = account.account-number
and
    account.branch-name = 'Perryridge')
```

- (Schema used in this example)





# Views


- Provide a mechanism to hide certain data from the view of certain users. To create a view we use the command:

**create view v as** <query expression>

where:

☞ <query expression> is any legal expression

☞ The view name is represented by v






# Example Queries

- A view consisting of branches and their customers

```
create view all-customer as  
  (select branch-name, customer-name  
   from depositor, account  
   where depositor.account-number = account.account-number)  
  union  
  (select branch-name, customer-name  
   from borrower, loan  
   where borrower.loan-number = loan.loan-number)
```

- Find all customers of the Perryridge branch

```
select customer-name  
  from all-customer  
  where branch-name = 'Perryridge'
```






# Derived Relations

- Find the average account balance of those branches where the average account balance is greater than \$1200.

```
select branch-name, avg-balance
from (select branch-name, avg (balance)
      from account
      group by branch-name)
      as result (branch-name, avg-balance)
where avg-balance > 1200
```

Note that we do not need to use the **having** clause, since we compute the temporary (view) relation *result* in the **from** clause, and the attributes of *result* can be used directly in the **where** clause.





# Modification of the Database – Deletion


- Delete all account records at the Perryridge branch

```
delete from account  
where branch-name = 'Perryridge'
```

- Delete all accounts at every branch located in Needham city.

```
delete from account  
where branch-name in (select branch-name  
                        from branch  
                        where branch-city = 'Needham')  
  
delete from depositor  
where account-number in  
      (select account-number  
      from branch, account  
      where branch-city = 'Needham'  
            and branch.branch-name = account.branch-name)
```

- (Schema used in this example)





# Example Query


- Delete the record of all accounts with balances below the average at the bank.

```
delete from account  
  where balance < (select avg (balance)  
    from account)
```

👉 Problem: as we delete tuples from *deposit*, the average balance changes

👉 Solution used in SQL:

1. First, compute **avg** balance and find all tuples to delete
2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)







# Modification of the Database – Insertion

- Add a new tuple to *account*

**insert into** *account*

**values** ('A-9732', 'Perryridge', 1200)

or equivalently


**insert into** *account* (*branch-name*, *balance*, *account-number*)

**values** ('Perryridge', 1200, 'A-9732')

- Add a new tuple to *account* with *balance* set to null

**insert into** *account*

**values** ('A-777', 'Perryridge', *null*)





# Modification of the Database – Insertion

- Provide as a gift for all loan customers of the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account

**insert into** *account*

**select** *loan-number, branch-name, 200*

**from** *loan*

**where** *branch-name* = 'Perryridge'

**insert into** *depositor*

**select** *customer-name, loan-number*

**from** *loan, borrower*

**where** *branch-name* = 'Perryridge'

**and** *loan.account-number = borrower.account-number*

- The select from where statement is fully evaluated before any of its results are inserted into the relation (otherwise queries like

**insert into** *table1 select \* from table1*

would cause problems





# Modification of the Database – Updates

- Increase all accounts with balances over \$10,000 by 6%, all other accounts receive 5%.


☞ Write two **update** statements:

```
update account  
set balance = balance * 1.06  
where balance > 10000
```

```
update account  
set balance = balance * 1.05  
where balance ≤ 10000
```

☞ The order is important

☞ Can be done better using the **case** statement (next slide)





# Case Statement for Conditional Updates

- Same query as before: Increase all accounts with balances over \$10,000 by 6%, all other accounts receive 5%.


**update** *account*

**set** *balance* = **case**

**when** *balance* <= 10000 **then** *balance* \* 1.05

**else** *balance* \* 1.06

**end**





# Update of a View

- Create a view of all loan data in *loan* relation, hiding the *amount* attribute

```
create view branch-loan as  
    select branch-name, loan-number  
    from loan
```

- Add a new tuple to *branch-loan*

```
insert into branch-loan  
    values ('Perryridge', 'L-307')
```

This insertion must be represented by the insertion of the tuple  
(*'L-307'*, *'Perryridge'*, *null*)

into the *loan* relation

- Updates on more complex views are difficult or impossible to translate, and hence are disallowed.
- Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation






# Joined Relations

- Join operations take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- Join condition – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- Join type – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

Join Types
<b>inner join</b>
<b>left outer join</b>
<b>right outer join</b>
<b>full outer join</b>

Join Conditions
<b>natural</b>
<b>on</b> <predicate>
<b>using</b> ( $A_1, A_2, \dots, A_n$ )





# Joined Relations – Datasets for Examples


## ■ Relation *loan*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

## ■ Relation *borrower*

<i>customer-name</i>	<i>loan-number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

- Note: borrower information missing for L-260 and loan information missing for L-155





# Joined Relations – Examples

- **loan inner join borrower on**  
*loan.loan-number = borrower.loan-number*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>	<i>loan-number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230

- **loan left outer join borrower on**  
*loan.loan-number = borrower.loan-number*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>	<i>loan-number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	<i>null</i>	<i>null</i>








# Joined Relations – Examples

## ■ *loan natural inner join borrower*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

## ■ *loan natural right outer join borrower*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	null	null	Hayes






# Joined Relations – Examples

- *loan full outer join borrower using (loan-number)*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	<i>null</i>	<i>null</i>	Hayes

- Find all customers who have either an account or a loan (but not both) at the bank.

**select** *customer-name*  
**from** (*depositor natural full outer join borrower*)  
**where** *account-number is null or loan-number is null*






# Data Definition Language (DDL)

Allows the specification of not only a set of relations but also information about each relation, including:


- The schema for each relation.
- The domain of values associated with each attribute.
- Integrity constraints
- The set of indices to be maintained for each relations.
- Security and authorization information for each relation.
- The physical storage structure of each relation on disk.





# Domain Types in SQL


- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *n* digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- Null values are allowed in all the domain types. Declaring an attribute to be **not null** prohibits null values for that attribute.
- **create domain** construct in SQL-92 creates user-defined domain types  
**create domain *person-name* char(20) not null**





# Date/Time Types in SQL (Cont.)

- **date.** Dates, containing a (4 digit) year, month and date
  - ☞ E.g. **date** '2001-7-27'
- **time.** Time of day, in hours, minutes and seconds.
  - ☞ E.g. **time** '09:00:30'      **time** '09:00:30.75'
- **timestamp:** date plus time of day
  - ☞ E.g. **timestamp** '2001-7-27 09:00:30.75'
- **Interval:** period of time
  - ☞ E.g. Interval '1' day
  - ☞ Subtracting a date/time/timestamp value from another gives an interval value
  - ☞ Interval values can be added to date/time/timestamp values
- Can extract values of individual fields from date/time/timestamp
  - ☞ E.g. **extract (year from r.starttime)**
- Can cast string types to date/time/timestamp
  - ☞ E.g. **cast** <string-valued-expression> **as date**





# Create Table Construct


- An SQL relation is defined using the **create table** command:

```
create table  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

- ☞  $r$  is the name of the relation
- ☞ each  $A_i$  is an attribute name in the schema of relation  $r$
- ☞  $D_i$  is the data type of values in the domain of attribute  $A_i$

- Example:

```
create table branch  
    (branch-name char(15) not null,  
    branch-city   char(30),  
    assets        integer)
```






# Integrity Constraints in Create Table

- **not null**
- **primary key** ( $A_1, \dots, A_n$ )
- **check** ( $P$ ), where  $P$  is a predicate

Example: Declare *branch-name* as the primary key for *branch* and ensure that the values of *assets* are non-negative.

```
create table branch  
    (branch-namechar(15),  
    branch-city  char(30)  
    assets       integer,  
    primary key (branch-name),  
    check (assets >= 0))
```

**primary key** declaration on an attribute automatically ensures **not null** in SQL-92 onwards, needs to be explicitly stated in SQL-89





# Drop and Alter Table Constructs

- The **drop table** command deletes all information about the dropped relation from the database.
- The **alter table** command is used to add attributes to an existing relation.

**alter table  $r$  add  $A$   $D$**

where  $A$  is the name of the attribute to be added to relation  $r$  and  $D$  is the domain of  $A$ .

☞ All tuples in the relation are assigned *null* as the value for the new attribute.

- The **alter table** command can also be used to drop attributes of a relation

**alter table  $r$  drop  $A$**

where  $A$  is the name of an attribute of relation  $r$

☞ Dropping of attributes not supported by many databases








# The *loan* and *borrower* Relations


<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

*loan*

<i>customer-name</i>	<i>loan-number</i>
Jones	L-170
Smith	L-230
Hayes	L-155


*borrower*





# The Result of *loan* inner join *borrower* on *loan.loan-number* = *borrower.loan-number*


<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>	<i>loan-number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230





# The Result of *loan* left outer join *borrower* on *loan-number*


<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>	<i>loan-number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	<i>null</i>	<i>null</i>





# The Result of *loan* natural inner join *borrower*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith





# Join Types and Join Conditions

## Join types

**inner join**

**left outer join**

**right outer join**


**full outer join**

## Join Conditions

**natural**

**on <predicate>**


**using  $(A_1, A_1, \dots, A_n)$**





# The Result of *loan* natural right outer join *borrower*


<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	<i>null</i>	<i>null</i>	Hayes





# The Result of *loan* full outer join *borrower* using(*loan-number*)

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	<i>null</i>	<i>null</i>	Hayes





# SQL Data Definition for Part of the Bank Database

**create table** *customer*

(*customer-name* **char**(20),  
*customer-street* **char**(30),  
*customer-city* **char**(30),  
**primary key** (*customer-name*))

**create table** *branch*

(*branch-name* **char**(15),  
*branch-city* **char**(30),  
*assets* **integer**,  
**primary key** (*branch-name*),  
**check** (*assets* >= 0))

**create table** *account*

(*account-number* **char**(10),  
*branch-name* **char**(15),  
*balance* **integer**,  
**primary key** (*account-number*),  
**check** (*balance* >= 0))

**create table** *depositor*

(*customer-name* **char**(20),  
*account-number* **char**(10),  
**primary key** (*customer-name*, *account-number*))

